

```
----
title: The Command R Model
slug: docs/command-r
hidden: false
description: >-
  Command R is a conversational model that excels in language tasks and supports multiple
  languages, making it ideal for coding use cases.
image: ../../../../assets/images/49841d1-cohere_meta_image.jpg
keywords: >-
  Cohere, large language models, generative AI, command model, chat models,
  conversational AI
createdAt: 'Tue Mar 05 2024 18:50:03 GMT+0000 (Coordinated Universal Time)'
updatedAt: 'Mon Jun 10 2024 14:22:50 GMT+0000 (Coordinated Universal Time)'
----
```

Command R is a large language model optimized for conversational interaction and long context tasks. It targets the “scalable” category of models that balance high performance with strong accuracy, enabling companies to move beyond proof of concept and into production.

Command R boasts high precision on [retrieval augmented generation](/docs/retrieval-augmented-generation-rag) (RAG) and tool use tasks, low latency and high throughput, a long 128,000-token context length, and strong capabilities across 10 key languages.

Model Details

Model Name	Description	
Context Length	Maximum Output Tokens	Endpoints
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		
----- ----- -----		

Command R August 2024 Release

Cohere's flagship text-generation models, Command R and Command R+, received a substantial update in August 2024. We chose to designate these models with time stamps, so in the API Command R 08-2024 is accesible with `command-r-08-2024`.

With the release, both models include the following feature improvements:

- For tool use, Command R and Command R+ have demonstrated improved decision-making around whether or not to use a tool.
- The updated models are better able to follow instructions included by the user in the preamble.
- Better structured data analysis for structured data manipulation.
- Improved robustness to non-semantic prompt changes like white space or new lines.
- Models will decline unanswerable questions and are now able to execute RAG workflows without citations

`command-r-08-2024` delivers around 50% higher throughput and 20% lower latencies as compared to the previous Command R version, while cutting the hardware footprint required to serve the model by half. Read more in the relevant blog post.

What's more, both these updated models can now operate in one of several safety modes, which gives developers more granular control over how models generate output in a variety of different contexts. Find more in these [safety modes docs] (<https://docs.cohere.com/docs/safety-modes>).

Unique Command R Model Capabilities

Command R has been trained on a massive corpus of diverse texts in multiple languages, and can perform a wide array of text-generation tasks. Moreover, Command R has been trained with a particular focus on excelling in some of the most critical business use-cases.

Multilingual Capabilities

We want Command R to serve as many people, organizations, and markets as possible, so the new Command R is capable of interacting in many languages to a fairly high degree of accuracy.

The model is optimized to perform well in the following languages: English, French, Spanish, Italian, German, Brazilian Portuguese, Japanese, Korean, Simplified Chinese, and Arabic.

Additionally, pre-training data has been included for the following 13 languages: Russian, Polish, Turkish, Vietnamese, Dutch, Czech, Indonesian, Ukrainian, Romanian, Greek, Hindi, Hebrew, Persian.

The model has been trained to respond in the language of the user. Here's an example:

```
```python PYTHON
co.chat(
 message="Écris une description de produit pour une voiture électrique en 50 à 75 mots"
)
```
```

And here's what the response might look like:

```
```text TEXT
Découvrez la voiture électrique qui va révolutionner votre façon de conduire. Avec son design élégant, cette voiture offre une expérience de conduite unique avec une accélération puissante et une autonomie impressionnante. Sa technologie avancée vous garantit une charge rapide et une fiabilité inégalée. Avec sa conception innovante et durable, cette voiture est parfaite pour les trajets urbains et les longues distances. Profitez d'une conduite silencieuse et vivez l'expérience de la voiture électrique!
```
```

Command R can not only be used to generate text in several languages but can also perform cross-lingual tasks such as translation or answering questions about content in other languages.

Retrieval Augmented Generation

Command R has been trained with the ability to ground its generations. This means that it can generate responses based on a list of supplied document snippets, and it will include citations in its response indicating the source of the information.

For more information, check out our dedicated guide on [retrieval augmented generation] (/docs/retrieval-augmented-generation-rag).

Tool Use

Command R has been trained with conversational tool use capabilities. This functionality takes a conversation as input (with an optional user-system preamble), along with a list of available tools. The model will then generate a json-formatted list of actions to execute on a subset of those tools. For more information, check out our dedicated [tool use] (/docs/tool-use) guide.